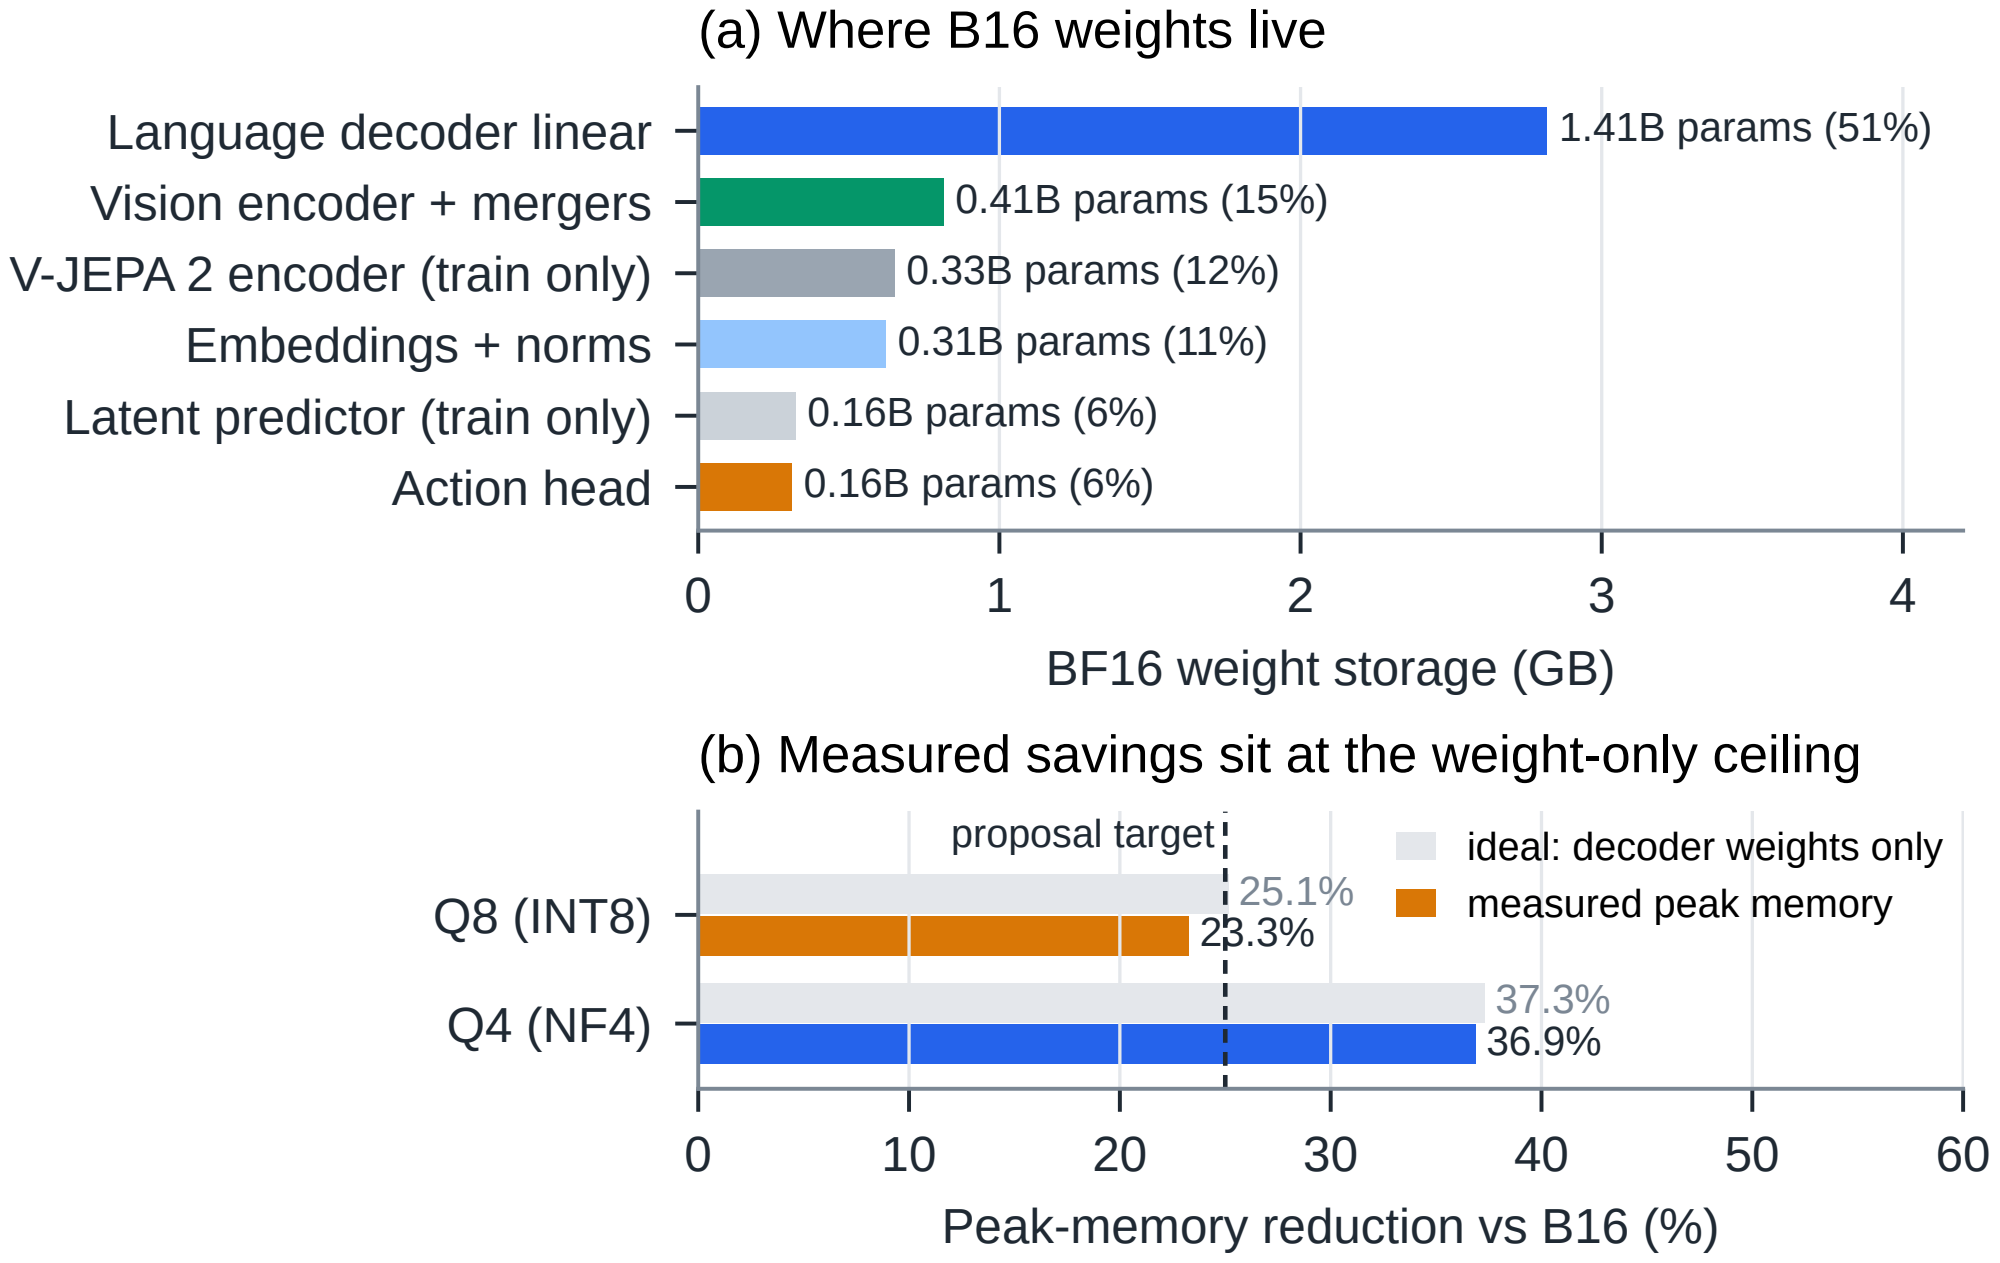
(a) Where B16 weights live
Language decoder linear
Vision encoder + mergers
V-JEPA 2 encoder (train only)
Embeddings + norms
Latent predictor (train only)
Action head
1.41B params (51%)
0.41B params (15%)
0.33B params (12%)
0.31B params (11%)
0.16B params (6%)
0.16B params (6%)
0
1
2
3
4
BF16 weight storage (GB)
(b) Measured savings sit at the weight-only ceiling
proposal target
ideal: decoder weights only
measured peak memory
Q8 (INT8)
Q4 (NF4)
25.1%
23.3%
37.3%
36.9%
0
10
20
30
40
50
60
Peak-memory reduction vs B16 (%)