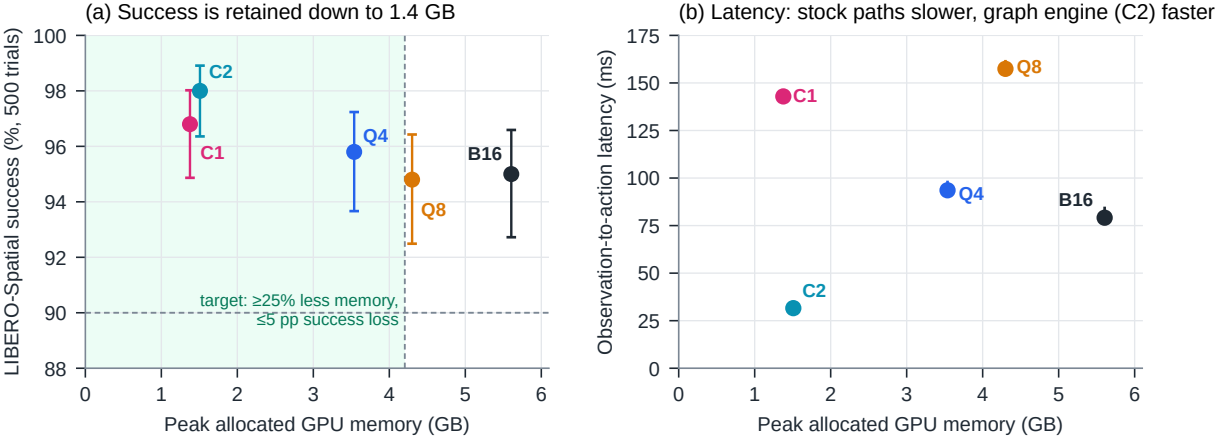
(a) Success is retained down to 1.4 GB
LIBERO-Spatial success (%, 500 trials)
Peak allocated GPU memory (GB)
C2
C1
Q4
Q8
B16
target: ≥25% less memory,
≤5 pp success loss
(b) Latency: stock paths slower, graph engine (C2) faster
Observation-to-action latency (ms)
Peak allocated GPU memory (GB)
Q8
C1
Q4
B16
C2