VLA-JEPA-LIBERO: confirmed compression tradeoffs

Released Qwen3-VL-2B policy · 500 reserved trials/configuration · RTX 3090 · no fine-tuning

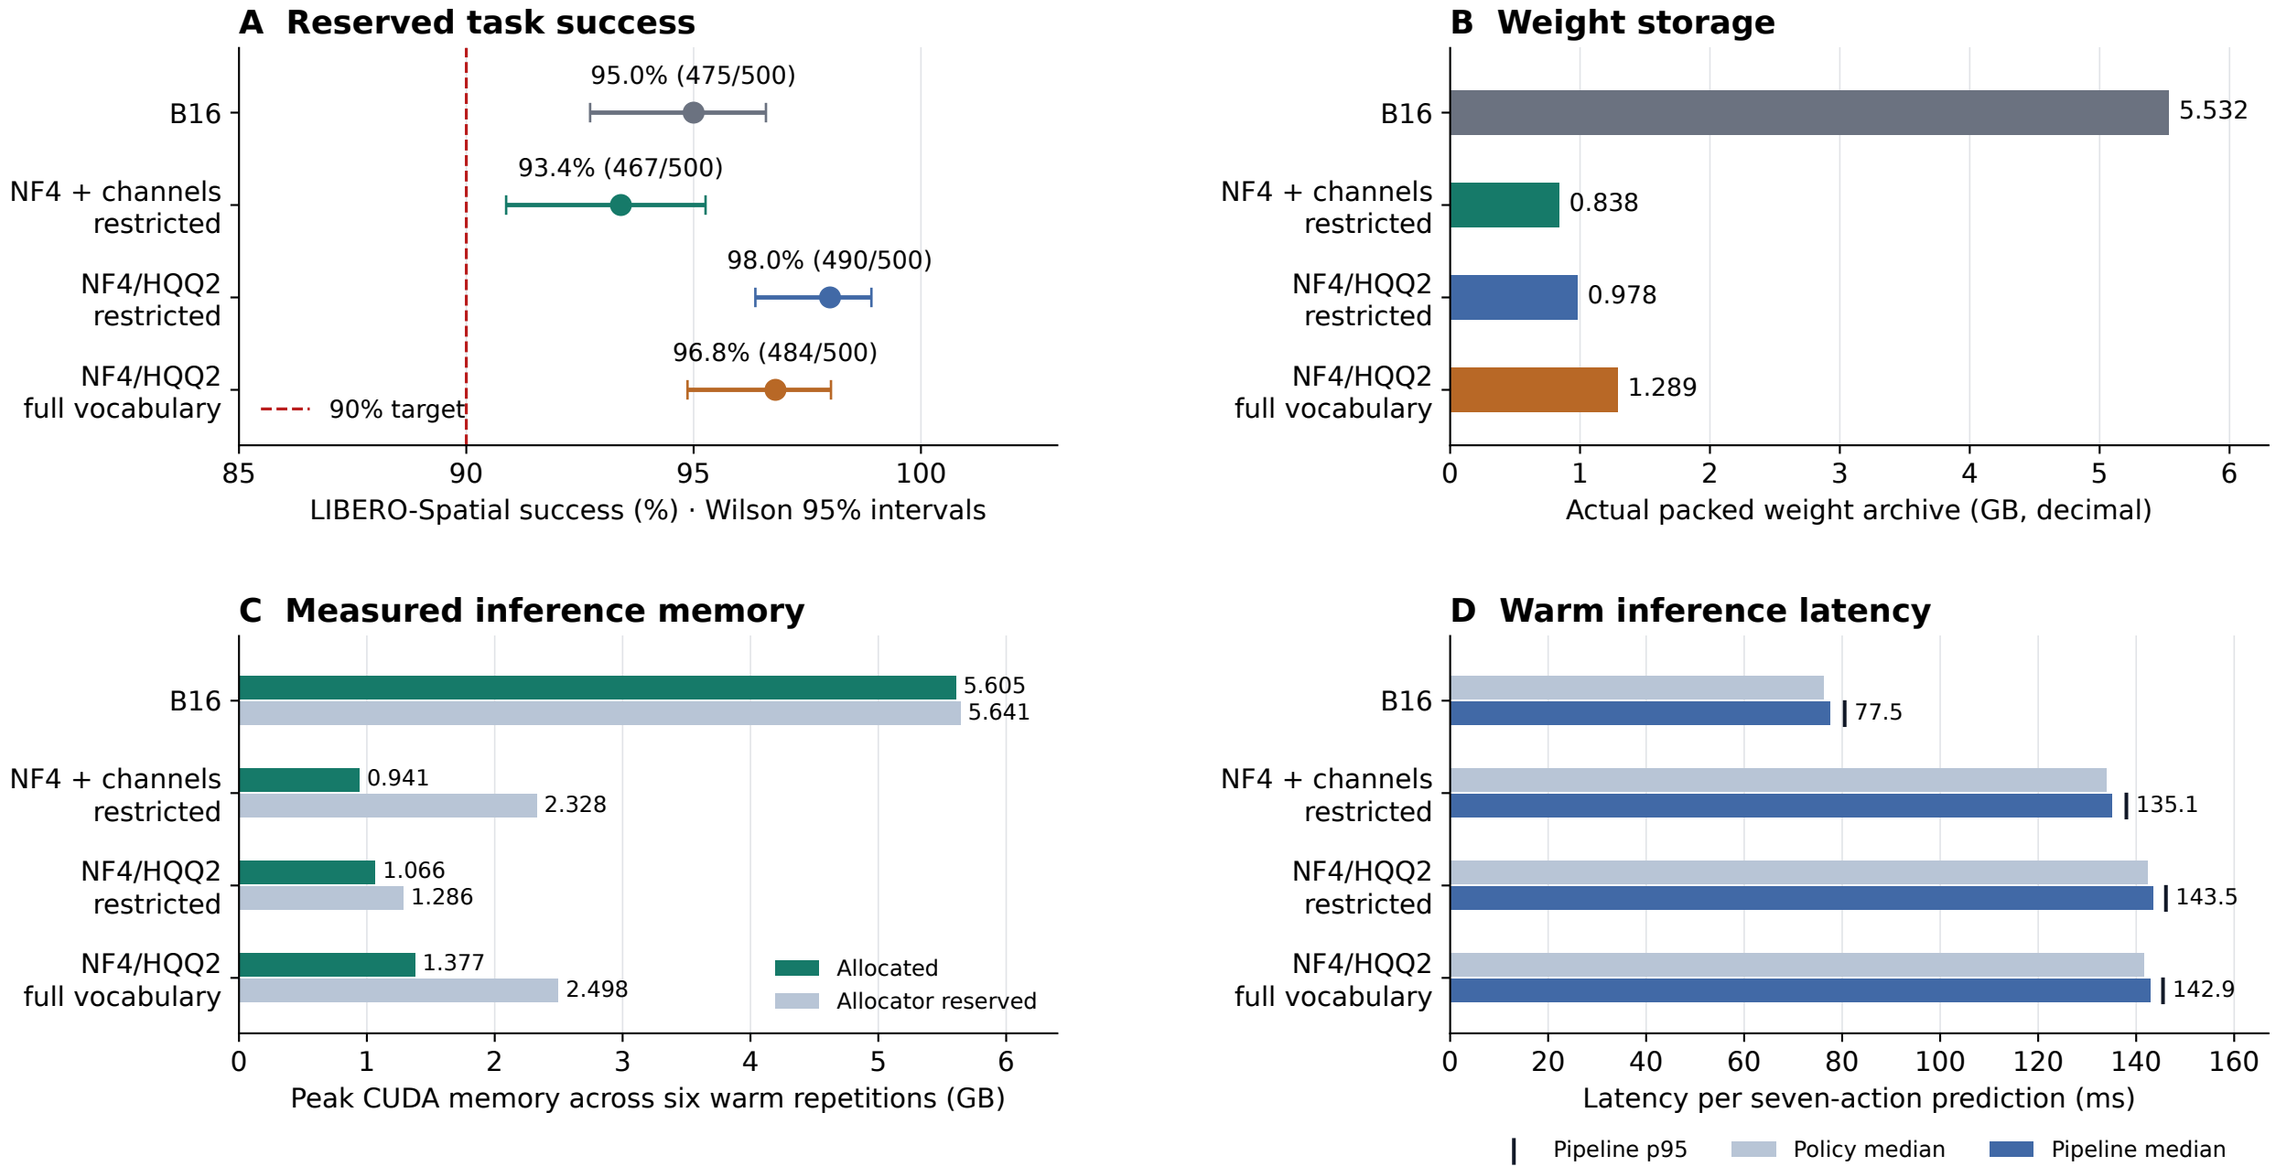

Timing: 500 predictions × 3 repetitions per mode; 20 warmups/repetition; shape priming separate.
Restricted deployments accept only 53 task-suite token IDs. GB = 10^9 bytes. CUDA memory is not NVML process memory.